

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("bmi.csv").dropna()
df
```

Out[3]:

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns

```
In [4]: df.head()
```

Out[4]:

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 500 entries, 0 to 499
Data columns (total 4 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Gender   500 non-null    object
 1   Height   500 non-null    int64
 2   Weight   500 non-null    int64
 3   Index     500 non-null    int64
dtypes: int64(3), object(1)
memory usage: 19.5+ KB
```

In [6]: `df.describe()`

Out[6]:

	Height	Weight	Index
count	500.000000	500.000000	500.000000
mean	169.944000	106.000000	3.748000
std	16.375261	32.382607	1.355053
min	140.000000	50.000000	0.000000
25%	156.000000	80.000000	3.000000
50%	170.500000	106.000000	4.000000
75%	184.000000	136.000000	5.000000
max	199.000000	160.000000	5.000000

In [7]: `df.columns`

Out[7]: `Index(['Gender', 'Height', 'Weight', 'Index'], dtype='object')`

In [8]: `feature_matrix = df[['Height', 'Weight']]`
`target_vector = df[['Index']]`

In [9]: `fs=StandardScaler().fit_transform(feature_matrix)`
`logr=LogisticRegression()`
`logr.fit(fs,target_vector)`

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 return f(*args, **kwargs)

Out[9]: `LogisticRegression()`

In [10]: `observation=[[1,2]]`

```
In [11]: prediction=logr.predict(observation)
         print(prediction)
```

```
[5]
```

```
In [12]:
         logr.classes_
```

```
Out[12]: array([0, 1, 2, 3, 4, 5], dtype=int64)
```

```
In [13]: logr.predict_proba(observation)[0][0]
```

```
Out[13]: 5.5956697582538237e-11
```

```
In [14]: logr.predict_proba(observation)[0][1]
```

```
Out[14]: 6.059900360819463e-10
```

Random Forest

```
In [15]: import numpy as np
         import pandas as pd
         import matplotlib.pyplot as plt
         import seaborn as sns
```

```
In [17]: df=pd.read_csv("bmi.csv")
```

```
In [18]: df['Gender'].value_counts()
```

```
Out[18]: Female    255
         Male      245
         Name: Gender, dtype: int64
```

```
In [19]: x=df[['Height','Weight']]
         y=df['Gender']
```

```
In [20]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [21]: from sklearn.ensemble import RandomForestClassifier
         rfc =RandomForestClassifier()
         rfc.fit(x_train,y_train)
```

```
Out[21]: RandomForestClassifier()
```

```
In [22]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
                    }
```

```
In [23]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc  
grid_search.fit(x_train,y_train)
```

```
Out[23]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                'min_samples_leaf': [5, 10, 15, 20, 25],  
                                'n_estimators': [10, 20, 30, 40, 50]},  
                    scoring='accuracy')
```

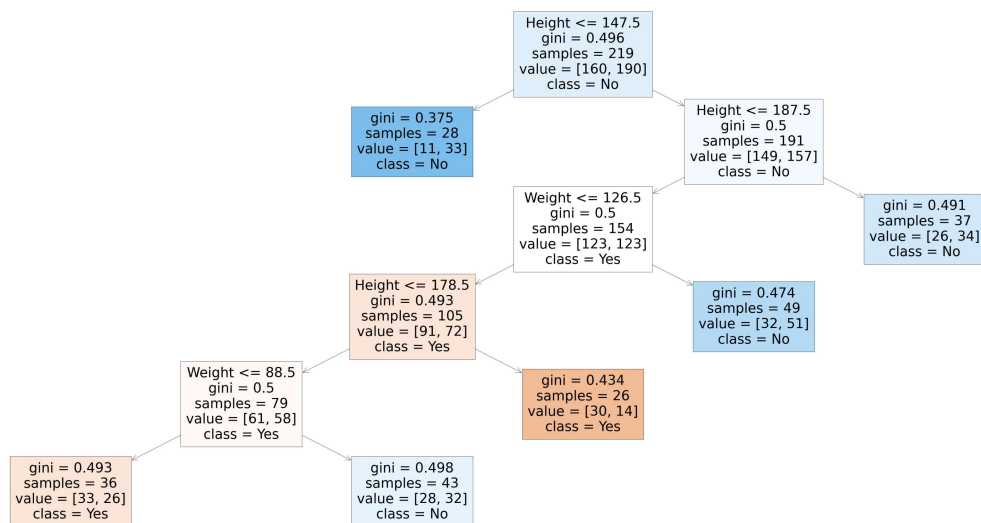
```
In [24]: grid_search.best_score_
```

```
Out[24]: 0.5485714285714286
```

```
In [25]: rfc_best=grid_search.best_estimator_
```

```
In [26]: from sklearn.tree import plot_tree
plt.figure(figsize=(89,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns, class_names=["Yes",
    )
```

```
Out[26]: [Text(2837.828571428571, 1993.2, 'Height <= 147.5\ngini = 0.496\nsamples = 219\nvalue = [160, 190]\nclass = No'),
Text(2128.3714285714286, 1630.8000000000002, 'gini = 0.375\nsamples = 28\nvalue = [11, 33]\nclass = No'),
Text(3547.2857142857138, 1630.8000000000002, 'Height <= 187.5\ngini = 0.5\nsamples = 191\nvalue = [149, 157]\nclass = No'),
Text(2837.828571428571, 1268.4, 'Weight <= 126.5\ngini = 0.5\nsamples = 154\nvalue = [123, 123]\nclass = Yes'),
Text(2128.3714285714286, 906.0, 'Height <= 178.5\ngini = 0.493\nsamples = 105\nvalue = [91, 72]\nclass = Yes'),
Text(1418.9142857142856, 543.5999999999999, 'Weight <= 88.5\ngini = 0.5\nsamples = 79\nvalue = [61, 58]\nclass = Yes'),
Text(709.4571428571428, 181.19999999999982, 'gini = 0.493\nsamples = 36\nvalue = [33, 26]\nclass = Yes'),
Text(2128.3714285714286, 181.19999999999982, 'gini = 0.498\nsamples = 43\nvalue = [28, 32]\nclass = No'),
Text(2837.828571428571, 543.5999999999999, 'gini = 0.434\nsamples = 26\nvalue = [30, 14]\nclass = Yes'),
Text(3547.2857142857138, 906.0, 'gini = 0.474\nsamples = 49\nvalue = [32, 51]\nclass = No'),
Text(4256.742857142857, 1268.4, 'gini = 0.491\nsamples = 37\nvalue = [26, 34]\nclass = No')]
```



In []: