

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("framingham.csv").dropna()
df
```

Out[3]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0
...
4231	1	58	3.0	0	0.0	0.0	0	1
4232	1	68	1.0	0	0.0	0.0	0	1
4233	1	50	1.0	1	1.0	0.0	0	1
4234	1	51	3.0	1	43.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0

3656 rows × 16 columns



```
In [4]: df.head()
```

Out[4]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	



In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3656 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  3656 non-null   int64
1   age                   3656 non-null   int64
2   education              3656 non-null   float64
3   currentSmoker          3656 non-null   int64
4   cigsPerDay             3656 non-null   float64
5   BPMeds                 3656 non-null   float64
6   prevalentStroke         3656 non-null   int64
7   prevalentHyp            3656 non-null   int64
8   diabetes               3656 non-null   int64
9   totChol                3656 non-null   float64
10  sysBP                  3656 non-null   float64
11  diaBP                  3656 non-null   float64
12  BMI                    3656 non-null   float64
13  heartRate              3656 non-null   float64
14  glucose                 3656 non-null   float64
15  TenYearCHD             3656 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 485.6 KB
```

In [6]: df.describe()

Out[6]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevaler
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656
mean	0.443654	49.557440	1.979759	0.489059	9.022155	0.030361	0
std	0.496883	8.561133	1.022657	0.499949	11.918869	0.171602	0
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1

In [7]: df.columns

Out[7]: Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds', 'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'], dtype='object')

```
In [8]: feature_matrix = df.iloc[:,0:15]
target_vector = df.iloc[:, -1]
```

```
In [9]: fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression()
logr.fit(fs,target_vector)
```

```
Out[9]: LogisticRegression()
```

```
In [10]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]]
```

```
In [11]: prediction=logr.predict(observation)
print(prediction)
```

```
[1]
```

```
In [12]:
logr.classes_
```

```
Out[12]: array([0, 1], dtype=int64)
```

```
In [13]: logr.predict_proba(observation)[0][0]
```

```
Out[13]: 0.0002214783507201723
```

```
In [14]: logr.predict_proba(observation)[0][1]
```

```
Out[14]: 0.9997785216492798
```

Random Forest

```
In [15]: df['TenYearCHD'].value_counts()
```

```
Out[15]: 0    3099
         1     557
         Name: TenYearCHD, dtype: int64
```

```
In [16]: x=df[['TenYearCHD']]
         y=df['TenYearCHD']
```

```
In [17]: g1={'TenYearCHD':{'0':1, "1":2}}
df=df.replace(g1)
df
```

Out[17]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0
...
4231	1	58	3.0	0	0.0	0.0	0	1
4232	1	68	1.0	0	0.0	0.0	0	1
4233	1	50	1.0	1	1.0	0.0	0	1
4234	1	51	3.0	1	43.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0

3656 rows × 16 columns



```
In [18]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [19]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[19]: RandomForestClassifier()

```
In [20]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]}
}
```

```
In [*]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc")
grid_search.fit(x_train,y_train)
```

In []: